

PREMISE SELECTION IS NOT ENOUGH: ACTION-CONDITIONAL EVIDENCE ALLOCATION FOR LEAN

Anonymous authors

Paper under double-blind review

ABSTRACT

Premise selection is usually treated as the problem of ranking which lemmas to send to a Lean prover. This view is incomplete: after names are selected, they must still be compiled into the proof interface that can consume them. We study *action-conditional evidence allocation*, where a method chooses not only selected names but also whether they enter Lean as Hammer facts, HammerCore inputs, simp lemmas, Aesop facts, Aesop simp rules, or elimination facts. We introduce ACE, a verified typed evidence compiler evaluated in a Mathlib 4.30 + LeanHammer 4.30 pre-theorem replay harness. On 230 replayable traced-corpus goals, the typed action grid reaches 58/230 kernel-verified proofs versus 38/230 for the best single action; a four-retry fixed portfolio, used as the strongest matched Lean-call control, reaches 57/230 out of fold and 58/230 when train-fitted. Exact Aesop controls sharpen the mechanism: oracle-core plus retrieved joint exposure solves 38/230 goals, but identity exposure solves 36/230, simps-only solves 35/230, retrieved-only joint exposure solves 34/230, and oracle-core-only joint exposure solves 35/230. Thus the robust claim is not a pure facts-versus-simps complementarity; it is that selected evidence has tactic-interface and source-dependent behavior that fixed typed portfolios can exploit. Stronger logistic and ComplementNB allocators do not beat or compress the fixed K=4 control, so we present adaptive routing as a boundary rather than a claim.

1 INTRODUCTION

Lean hammers and neural theorem provers depend on premise selection: given a proof goal, the system must choose accessible lemmas that make proof search tractable and reconstructable in Lean. Modern systems therefore focus on retrieving better premises before running a backend prover (Yang et al., 2023; Zhu et al., 2025; Gao et al., 2026). This framing hides a second decision that is just as important in Lean: once a name is selected, it must be assigned to a concrete proof interface. A name can be used as a Hammer premise, a HammerCore fact, a simp lemma, an Aesop fact rule, an Aesop simp rule, or part of a tactic-specific local context. These interfaces are not interchangeable, so a retrieved premise list is not yet an executable proof-search intervention.

The central challenge is that useful evidence is not monotonic under either quantity or interface. Adding more names can improve proof-core coverage, but it can also increase search branching, timeout risk, and reconstruction ambiguity. More subtly, exposing the same source of evidence through the wrong tactic interface can make it inert or harmful even when the names are relevant. Our Mathlib mechanism-isolation experiments make this visible using a mixed source pool of traced oracle-core names and learned retrieved candidates. Aesop with this mixed source exposed jointly as facts and simp lemmas proves 38/230 replayable traced-corpus goals, compared with 29/230 for empty Aesop. However, exact controls show that most of this lift is not uniquely caused by facts+simps complementarity: identity exposure proves 36/230, simps-only proves 35/230, facts-only proves 32/230, retrieved-only joint exposure proves 34/230, and oracle-core-only joint exposure proves 35/230. The key observation is therefore not simply that “both Aesop channels are necessary”, but that selected evidence behaves differently under source choice, channel assignment, and Lean search budget.

We propose ACE, an action-conditional evidence allocator for bounded Lean proof search. Given a replayable theorem context and a candidate set, ACE compiles selected names into typed proof

actions: Hammer facts, HammerCore inputs, simplification lemmas, `solve_by_elim` facts, and Aesop fact/simp channels. The method is deliberately compute-budgeted: every comparison starts from the same empty-premise Hammer attempt and then spends a fixed number of retry actions. This makes the main question reviewer-checkable: how much of the verified oracle headroom comes from assigning evidence to the right interface, and can any learned allocator beat a strong fixed typed control?

The main verified result is on non-generated replayable Mathlib theorem contexts. We migrate traced goals to a Mathlib 4.30 + LeanHammer 4.30 pre-theorem replay harness and identify 230 replayable goals from 490 cleaned traces. On this subset, the typed action grid reaches a 58/230 oracle, compared with 38/230 for the best single action. A fixed portfolio of four typed retry actions after `hammer_empty` reaches 57/230 out of fold and 58/230 when train-fitted. This fixed sequence is not the proposed novelty; it is the hard matched Lean-call control that shows how large the interface-typing effect is. Stronger logistic and ComplementNB allocators match but do not beat the $K=4$ fixed control, so adaptive routing remains a future target.

This paper makes three contributions. First, we formulate action-conditional evidence allocation as a distinct problem from premise retrieval: selected evidence must be compiled into a Lean interface, not only ranked. Second, we implement and evaluate ACE in a real Mathlib 4.30 + LeanHammer 4.30 replay harness, showing that small typed action sets recover nearly all current oracle headroom on the replayable subset. Third, we provide mechanism evidence and boundaries: Aesop is sensitive to source and exposure choice, broad 32-name insertion is harmful in the original matrix, strict syntactic filtering does not explain the gains, and learned allocators do not yet beat the fixed typed control. Trace-core failure-feedback experiments remain important discovery evidence, but the current supported main claim is action-conditional evidence allocation.

The scope is deliberately narrower than a full theorem-proving leaderboard. We treat fixed typed portfolios as strong matched Lean-call controls, not as weak baselines, and we do not claim adaptive routing until it beats them. We also distinguish controlled lexical retrieval proxies from full LeanSearch v2 reproductions, and bridge replay from open-ended proof generation. These boundaries keep the paper’s claim reviewer-checkable: the contribution is a typed interface mechanism for selected evidence in Lean proof search.

2 RELATED WORK

Premise selection for Lean. Premise selection is a long-standing interface between theorem provers and large formal libraries. LeanDojo introduced infrastructure and benchmarks for retrieval-augmented theorem proving in Lean (Yang et al., 2023). Recent work on Lean hammers studies learned premise ranking and ATP integration for Lean (Zhu et al., 2025), while LeanSearch v2 focuses on global premise retrieval in Lean 4 (Gao et al., 2026). These works motivate our setting and provide the strongest conceptual baselines: retrieve useful premises from a large library before proof search. Our distinction is that ACE studies the next interface decision after retrieval. It asks how selected names should be compiled into Hammer facts, HammerCore inputs, simp lemmas, Aesop fact rules, Aesop simp rules, or tactic-specific local evidence.

Tactic selection and tactic portfolios. TacticToe learns tactic-level proof search in HOL4 and combines tactic prediction with premise selection inside a search procedure (Gauthier et al., 2021). This is the closest conceptual objection to our framing: a typed action grid can look like another tactic portfolio. Our object is different. Prior tactic-selection systems choose which tactic or tactic sequence to try; ACE holds the selected evidence fixed enough to study how the same names should be represented for the consuming interface, including multi-channel assignment inside a single Aesop invocation. The ablations therefore compare fact-only, simp-only, joint, widened, and source-decomposed evidence programs, not only different tactic names.

Proof repair, compiler feedback, and agentic provers. Several recent systems use Lean feedback to repair or generate proof text. APRIL learns to repair Lean proofs from compiler feedback (Wang et al., 2026); APOLLO and OProver frame formal proving as a broader agentic collaboration or whole-proof search problem (Ospanov et al., 2025; Ma et al., 2026). Process-verified reinforcement learning also uses Lean verification as a training signal (Kim & Yun, 2026). ACE is complementary

to these systems: it does not edit proof text or search over natural-language proof plans. Its object is the typed action by which selected names are supplied to Lean’s existing proof mechanisms.

Search guidance and local context. Search-guidance methods such as LeanProgress predict proof progress to steer theorem proving (Huang et al., 2025), while long-context theorem proving benchmarks such as miniCTX emphasize the importance of local and user-provided context (Hu et al., 2024). These lines share our view that static global retrieval is incomplete. Our experiments focus on a different mechanism: typed exposure of selected evidence under fixed retry budgets. Failure feedback remains useful in our trace-core discovery experiments, but the verified Mathlib replay result shows that interface choice itself can dominate low-capacity adaptive routing.

3 PROBLEM FORMULATION

Action-conditional evidence allocation under a fixed retry budget. Let g be a Lean goal and let $\mathcal{A}(g)$ be the accessible name set induced by imports, namespaces, declarations, and local context. A typed proof action is a pair $a = (\tau, P)$, where τ is a Lean interface such as Hammer, HammerCore, simplification, `solve_by_elim`, or Aesop, and $P \subset \mathcal{A}(g)$ is the set of names exposed through that interface. More generally, a typed evidence program is an assignment

$$z_{p,\tau} \in \{0, 1\}, \quad \sum_{p \in \mathcal{A}(g)} z_{p,\tau} \leq b_\tau,$$

where $z_{p,\tau} = 1$ means that name p is exposed to interface channel τ under channel budget b_τ . The same name may receive multiple interface labels: an Aesop joint facts+sims action is one compiled program, not two independent retries. A compiler $C(z)$ maps this assignment to Lean tactic syntax, such as Hammer facts, HammerCore inputs, simplifier lemmas, `solve_by_elim` facts, Aesop safe facts, and Aesop simp rules. This separates four decisions that are often conflated: candidate provenance, channel eligibility, channel assignment, and outer action scheduling.

An evidence allocator chooses a sequence of typed programs,

$$a_t = (\tau_t, P_t), \quad |P_t| \leq k_t,$$

and runs each action under a search budget B_t . The prover returns either a kernel-verified success or a failure observation F_t , such as a timeout, missing premise signal, type mismatch, typeclass failure, failed simplification, or reconstruction failure. A closed-loop selector is a policy

$$a_{t+1} = \pi(g, a_{\leq t}, F_t, H_t),$$

where H_t includes previous attempt history and candidate features. The objective is to maximize verified success under a fixed total retry budget, rather than to maximize static premise recall at a single k . A fixed typed portfolio is one allocator; a learned action router is another. The experiments ask whether interface typing itself creates headroom before claiming that learned routing has solved the allocation problem.

Evaluation axes. We separate three notions that are often conflated. *Kernel-verified proof-action success* means that a patched Lean file containing a candidate proof action checks successfully. This is the main verified metric for ACE. *Trace-core success* means that the selected premise set covers the traced proof-core premises for a goal; it is a premise-recovery metric, not a full proving metric. *Bridge verified success* is stricter than trace-core success: trace-core recovery must coincide with successful replay of the original traced tactic script in a real Lean process. Bridge verification is still not a full proof-generation benchmark, because it replays known tactic scripts, but it filters out trace-core cases that drift too far from Lean execution.

4 METHOD

4.1 OVERVIEW

ACE treats a selected name as incomplete until it is assigned to a proof-action interface. Figure 1 gives the full loop. The system first builds a candidate pool and runs a shared empty-premise Hammer




Figure 1: ACE evaluates action-conditional evidence allocation under a fixed retry budget. Selected names are compiled into typed Lean interfaces, and each routed action is checked by the Lean kernel.

attempt. It then spends a fixed retry budget on typed actions that expose selected names as Hammer facts, HammerCore inputs, simp lemmas, solve_by_elim facts, or Aesop facts/simp rules. The important design choice is that the retry budget is spent on compiling evidence into interfaces, not on blindly expanding a single untyped premise list.

4.2 TYPED EVIDENCE PROGRAMS

The allocator is built from action families that consume selected names differently. Hammer actions pass names as ordinary premises; HammerCore actions expose a smaller internal proof-search interface; simplification actions use selected names as rewrite/simp lemmas; solve_by_elim uses names as local facts; Aesop actions can add names as fact rules, simp rules, or both. For each replayable theorem context, we materialize a patched pre-theorem Lean file for each action and classify the result by running Lean with a fixed timeout.

The fixed portfolio is chosen greedily on the training split and evaluated out of fold. Every policy receives the same first action, hammer_empty, and then a retry budget K . This protocol makes the strongest control explicit: an adaptive policy is only useful if it beats or compresses the best fixed sequence of typed actions under the same Lean-call budget. In the current matrix, the train-fitted four-action portfolio after hammer_empty is aesop_core_plus_learned, hammerCore_core_plus_learned, hammer_core_plus_learned16, and solve_by_elim_core.

This design also prevents a common misreading of the method. The fixed portfolio is not intended to prove that a hand-written order is optimal; it is the hard control that any learned action router must beat. Its near-oracle performance is therefore evidence that interface typing itself is a large effect, while the absence of an adaptive win is reported as a boundary rather than hidden as a negative result.

4.3 AESOP FACT/SIMP EXPOSURE

Aesop is the most important new interface in the verified matrix, so we isolate how selected names enter it. Given a ranked name list, we split available names into fact candidates and simp candidates, then instantiate Aesop in three exposure modes: facts-only, simps-only, and facts+simps. We also vary source pools: traced oracle-core names, retrieved-only names, and oracle-core plus retrieved names at budgets 8, 16, and 32. This gives a matched source/budget control: the goal and candidate source are fixed, while the channel assignment is removed, added, or widened.

The Aesop ablation is not a new oracle-improving action set; its role is explanatory. The original broad ablation suggested a strong facts+simps interaction, but exact matched controls give a more conservative readout. For the best current oracle-core plus retrieved source, joint exposure recovers the 38/230 Aesop result, yet identity exposure reaches 36/230, simps-only reaches 35/230, facts-only reaches 32/230, and joint-only over exact single-channel controls is only 2 goals. Retrieved-only joint exposure still reaches 34/230 and oracle-core-only reaches 35/230. The supported mechanism is therefore source/exposure/search sensitivity, not a clean causal claim that facts and simps are complementary in isolation.

Table 1: Main claim-evidence map. The primary evidence is the verified typed action matrix and Aesop channel counterfactual; trace-core and bridge rows are supporting discovery evidence.

Claim	Strong control	Supported result	Readout
Typed actions matter	Best single action: 38/230	Oracle typed grid: 58/230	+20 goals
Small portfolios are hard controls	Empty Hammer: 29/230	Fixed K=4 OOF: 57/230	near oracle
Aesop exposure is source/interface sensitive	Empty Aesop: 29/230; identity: 36/230	Oracle-core+retrieved joint: 38/230 ; retrieved-only: 34/230	modest joint/source effect
Broad insertion can hurt	Oracle-core+retrieved8 facts+sims: 38/230	Oracle-core+retrieved32 facts+sims: 3/230 in original matrix	negative boundary
Adaptive allocation is not yet supported	Fixed K=4 OOF: 57/230	Best learned allocator: 57/230 at K=4	boundary
Failure feedback is useful discovery evidence	Learned+base fallback: 86.3% mean	Trace-core final-base8: 92.9% mean	supporting result
Bridge replay supports premise recovery	Fallback bridge: 44.5%	Final-base8 bridge: 53.5%	supporting result

4.4 TRACE-CORE DISCOVERY CONTROLLERS

The typed-portfolio study was motivated by earlier trace-core controllers. In local Mathlib settings, FAR-HAMMER uses visible candidate features and failure types to decide how to rerank or shrink premise sets. In imported-core settings, where lexical and same-file priors are too weak, FAR-HAMMER trains a global retriever from traced proof-core labels and then trains a second-stage scorer conditional on the first failure. These controllers are not the verified main method in this draft, but they provide discovery evidence for two ideas that survive in the verified study: premise budgets are non-monotonic, and retry decisions should be evaluated against strong fixed-budget controls.

4.5 WHY TYPED ACTIONS SHOULD HELP

Let $Y_g(z; B) \in \{0, 1\}$ denote the kernel-verified potential outcome for goal g when typed assignment z is checked under search budget B . The marginal intervention effect of exposing name p through interface τ is

$$\Delta_g(p, \tau \mid z, B) = Y_g(z + e_{p, \tau}; B) - Y_g(z; B).$$

This formulation is intentionally intervention-level rather than only a marginal retrieval score. It permits the same name to have positive effect in one interface and negative effect in another, allows facts and simp rules to interact through the current program z , and permits larger evidence sets to have negative marginal effect under bounded search. A typed action helps when the compiler finds assignments whose realized intervention effects are positive under the Lean backend’s budget, not merely when a retriever ranks relevant names highly.

5 EXPERIMENTS

5.1 PROTOCOL

We evaluate two evidence layers. The main verified layer is a Mathlib 4.30 + LeanHammer 4.30 pre-theorem replay harness. We start from 490 cleaned traced-corpus goals, keep the 230 goals whose original traced tactic script replays in the patched pre-theorem context, and then run typed proof actions in Lean. Success means that the patched Lean file checks successfully.

The second layer is trace-core discovery evidence from earlier premise-selection experiments. There, success means that the selected premises cover the proof-core premises extracted from traced proofs. We keep these results because they explain why premise budgets and failure feedback matter, but they are not the main verified claim. BM25 and iterative lexical rows are controlled retrieval policies inside our candidate-pool protocol, not full system-to-system reproductions of the published LeanSearch v2 pipeline.

Table 2: Verified typed proof-action matrix on 230 replayable Mathlib 4.30 traced-corpus theorem contexts. Success means the patched Lean file checks.

Matrix	Attempts	Verified attempts	Best single	Oracle
Original typed grid	3450	246	31/230	51/230
Extended typed grid	7130	506	38/230	58/230
Full Aesop-ablation grid	11500	582	38/230	58/230

Table 3: Budgeted allocator gate after `hammer_empty`. Learned allocators match but do not beat or compress the fixed K=4 typed control.

Retry budget K	Fixed OOF	Pure logreg	Residual logreg	Residual CNB	Oracle
1	49/230	41/230	41/230	38/230	58/230
2	55/230	41/230	55/230	52/230	58/230
3	55/230	44/230	56/230	55/230	58/230
4	57/230	47/230	57/230	57/230	58/230

5.2 VERIFIED TYPED PROOF-ACTION MATRIX

Table 2 reports the main verified experiment. The original scaled matrix already shows action-dependent headroom: 51/230 oracle versus 31/230 for the best single action. Extending the action set with Aesop and additional typed interfaces raises the oracle to 58/230 and the best single action to 38/230. The full Aesop-ablation matrix does not raise the oracle further, which means the ablation is explanatory rather than a ceiling improvement.

The strict action-dependent subset contains 29 goals where `hammer_empty` fails but some non-empty or non-default action succeeds. Aesop solves 20/29 strict goals and is the only successful family for 6, HammerCore solves 12/29 and is the only family for 5, and `solve_by_elim` contributes 5/29 with 2 only-family cases. This distribution rules out a Hammer-only interpretation and motivates a typed-interface portfolio.

5.3 BUDGETED ALLOCATORS AND AESOP COUNTERFACTUALS

Table 3 evaluates fixed and learned allocators under matched retry budgets after the shared `hammer_empty` attempt. The fixed typed portfolio is the strongest current policy: it reaches 55/230 with two retries, 57/230 with four retries out of fold, and 58/230 when the four-action sequence is fit on all training data. Stronger text/numeric logistic and ComplementNB allocators, trained out of fold from goal text, first-failure status, and typed fact/simp pool features, do not beat or compress this control. Adaptive allocation is therefore an explicit future target, not a main claim.

Table 4 checks whether the K=4 result is simply repeated use of one family. It is not: four Aesop retries reach 49/230, HammerCore-only and simplification-only reach 41/230, and Hammer-only reaches 32/230, while the typed grid reaches 57/230 out of fold. Random K=4 portfolios over the full grid average 41.9/230 over 100 seeds. The fixed typed sequence is therefore a strong control because it combines interfaces, not because it spends more calls inside one backend.

Table 5 reports empirical runner wallclock rather than treating all calls as equal compute. These numbers are not heartbeat-normalized costs, so we use them as a secondary frontier. They support the conservative wording that the main comparisons are matched by Lean-call budget, with wallclock reported for transparency.

Table 6 isolates the strongest interface under the exact P0 controls. Joint exposure remains the best Aesop configuration, but the exact controls prevent an over-strong causal interpretation: identity exposure, simps-only exposure, and single-source joint exposure account for most of the lift over empty Aesop. The result is still useful because it shows that the evidence compiler must expose the source/channel choice explicitly, but it is no longer evidence for a large pure facts-versus-simps complementarity.

Table 4: Homogeneous K=4 controls under the same outer Lean-call budget after `hammer_empty`. Typed diversity is stronger than spending all retries inside one family.

Portfolio family	OOF K=4	Strict goals	Avg calls
Full typed grid	57/230	28/29	4.18
Aesop-only	49/230	20/29	4.23
HammerCore-only	41/230	12/29	4.34
Simplification-only	41/230	12/29	4.34
<code>solve_by_elim</code> -only	34/230	5/29	4.43
Hammer-only	32/230	3/29	4.47
Raw arithmetic/normalization	29/230	0/29	4.50

Table 5: First-success runner-wallclock frontier. The primary protocol is matched Lean calls; wallclock is reported as an empirical secondary cost.

Policy	Success	Avg calls	Avg time (s)
<code>hammer_empty</code>	29/230	1.00	8.31
<code>aesop_core_plus_learned</code> after empty	49/230	1.87	16.36
Full typed train-fitted K=2	55/230	2.66	26.37
Full typed train-fitted K=4	58/230	4.18	40.07

5.4 SUPPORTING TRACE-CORE DISCOVERY EVIDENCE

The trace-core and imported-core experiments are discovery evidence rather than the verified main claim. They explain why the typed-action protocol is shaped around bounded retries and hard controls: local trace-core recovery showed that failure-conditioned feature use can improve real replay bridges, timeout stress showed that larger premise sets can be worse than one-shot retrieval, and imported-core experiments showed that local heuristics fail when proof-core premises come from global library context. The strongest bridge-facing discovery numbers are 72.0% local bridge verified success for the failure-aware controller versus 43.0% for blind expansion, and 53.5% imported bridge verified success for final-base8 versus 44.5% for learned+base fallback. We keep these results as provenance for the current design, while the main claim remains the kernel-verified typed evidence compiler in the 230-goal replay harness. Appendix 8 summarizes the representative discovery numbers.

5.5 ABLATIONS AND FAILURE DIAGNOSIS

The ablations in Table 7 identify why the final method is shaped as typed evidence allocation rather than a larger untyped premise list or a learned router claim. The exact Aesop controls show a useful but modest source/exposure effect, while ruling out the earlier stronger interpretation that facts and_simps form a large isolated complementarity. Strict interface filtering and learned allocation gates are useful boundaries: they show that the result is not explained by simple syntactic cleanup, and that current learned routers do not yet beat the hard fixed control.

6 THEORY AND ANALYSIS

Proposition 1 (Typed interface interventions can change verified outcomes). *Let $Y_g(z; B) \in \{0, 1\}$ be the kernel-verified outcome for goal g under typed evidence assignment z and search budget B . If there exist a name p and two interfaces τ_1, τ_2 such that*

$$Y_g(z + e_{p, \tau_1}; B) \neq Y_g(z + e_{p, \tau_2}; B),$$

then no interface-agnostic premise score for p can represent the full intervention effect of exposing p to Lean.

Sketch. The intervention $e_{p, \tau}$ changes both the syntax sent to Lean and the backend search process that consumes the name. If two interface assignments yield different verified outcomes under the same goal and budget, a scalar premise ranking over p alone has lost information needed to choose the successful action. This is the formal reason for compiling selected evidence into typed proof actions rather than treating retrieval as the final decision.

Table 6: Exact Aesop source/channel controls. Joint oracle-core plus retrieved exposure is best, but matched controls show that the mechanism is source/exposure sensitivity rather than a large pure channel-complementarity effect.

Aesop control	Verified goals	Interpretation
empty Aesop	29/230	default baseline
retrieved-only, joint	34/230	deployable-source anchor is nonzero
oracle-core-only, joint	35/230	oracle source alone is similar
oracle-core+retrieved, joint	38/230	best exact Aesop row
oracle-core+retrieved, identity	36/230	identity explains much of the lift
oracle-core+retrieved, simps-only	35/230	simp exposure is strong
oracle-core+retrieved, facts-only	32/230	fact exposure is weaker but nonzero
oracle-core+retrieved, random split	32/230	channel assignment still matters modestly

Table 7: Focused diagnostic results. The main verified insight is action-conditional evidence allocation; trace-core rows provide supporting discovery evidence for non-monotonic premise budgets and failure feedback.

Question	Evidence	Interpretation
Does Aesop need both exposure channels?	Exact joint: 38/230; identity: 36/230; simps-only: 35/230; facts-only: 32/230	Strong complementarity is not supported; source/exposure sensitivity remains.
Does adding more Aesop names help?	Oracle-core+retrieved 8 facts+simps: 38/230; oracle-core+retrieved 32 facts+simps: 3/230	Broad insertion can hurt.
Is strict syntactic filtering enough?	Filtered-only oracle: 4/230; combined oracle remains 58/230; new oracle goals: 0	Cleanup alone is not the mechanism.
Does learned allocation beat fixed K=4?	Fixed K=4 OOF: 57/230; best learned gate: 57/230 at K=4	Adaptive routing is not a main claim.
Are visible features enough statically?	Static all-feature rerank: 86.5%; failure-conditioned: 95.7%	Failure event changes feature utility.
Which visible features matter most?	Name-only FAR-HAMMER: 95.0%; full visible FAR-HAMMER: 95.7%	Namespace/name morphology is strongest.
Is failure type alone enough?	Failure-type-only: 91.5%, equal to top-k expansion	Candidate-level conditioning is required.
Does broad base guardrail help?	Hard guardrails drop bridge trace from 72.0% to 63.0/60.0%	Broad rescue damages ranking.
Does expert mixing help?	Multi-expert bridge trace drops to 62.0%; expert-16 drops final-base 8	Hand-written expert mixing is unstable.
Are misses candidate-pool failures?	200-bridge inspection: 0/17 candidate-pool misses	Remaining headroom is ranking/calibration.

Proposition 2 (Premise coverage is not monotonic in premise budget). *There can exist premise sets $P_1 \subset P_2$ such that P_1 covers enough useful premises for the backend to succeed under budget B , while P_2 contains the same useful premises but causes timeout or reconstruction failure under the same budget.*

Sketch. Adding premises preserves proof-core coverage but can increase branching, generated clauses, simplification search, or reconstruction ambiguity. Under fixed time/search budget, the prover may fail on the larger set even though the smaller set contains the same proof-relevant lemmas. Tables 7 and 6 are empirical versions of this proposition: top- k expansion can hurt under timeout pressure, and top-32 Aesop exposure is worse than smaller typed exposure.

Proposition 3 (Complementary channels can dominate either channel alone). *Suppose an interface has two evidence channels τ_f and τ_s that enable different proof-search transitions, and a goal requires transitions from both channels within budget B . Then exposing P_f only through τ_f or P_s only through τ_s can fail, while the joint typed action $(\tau_f, P_f) \cup (\tau_s, P_s)$ succeeds under the same outer retry budget.*

Sketch. The claim is existential. A fact channel can introduce the relational or propositional facts needed to close branches, while a simp channel can normalize terms or rewrite hypotheses so that

those facts become applicable. If either channel is absent, search may fail to reach the closing state within budget; if both are present, the interface can interleave normalization and rule application. Table 6 is also a guardrail against overclaiming: in our exact controls, joint exposure is best at 38/230, but identity exposure reaches 36/230 and simpsons-only reaches 35/230, so the empirical mechanism is source/exposure sensitivity rather than a large isolated complementarity effect.

Proposition 4 (A guardrail should be narrow when the learned controller is already strong). *If a learned failure-conditioned scorer improves most retry decisions but occasionally suppresses high-base-score rescue premises, then a small late-stage base guardrail can improve bridge recovery while broad base mixing can reduce success.*

Sketch. The learned controller and the base retriever encode different biases. A broad mixture changes many rankings, including rankings already corrected by failure feedback. A narrow final guardrail only affects the tail cases where high-base premises are about to be excluded. This matches the trace-core discovery pattern: final-base8 improves bridge verified success, while hard guardrails and broad expert/base mixing are harmful. It also motivates the verified typed-action protocol, where broad insertion is treated as a hypothesis to test rather than as a default improvement.

7 LIMITATIONS

The main limitation is evaluation scope. The verified typed-action results are kernel-checked, but they are measured on the 230-goal replayable subset of a 490-goal cleaned trace pool, not on every Mathlib theorem and not as an open-ended theorem-proving benchmark. The harness patches original files in a pre-theorem context, which is the right setting for avoiding target-theorem leakage, but it also means non-replayable migrated traces are outside the current main evaluation.

A second limitation is that the current strongest policy is a fixed typed portfolio, not a learned adaptive router. The budgeted policy experiments are useful because they prevent overclaiming: NB/kNN, logistic, and ComplementNB allocators do not beat or compress the fixed typed portfolio under matched Lean-call budget. The trace-core experiments use supervised proof-core traces and should be read as discovery evidence for premise budgets and failure feedback, not as the main verified claim. Future work should either train a stronger adaptive allocator that beats the fixed typed control or expand the replayable verified corpus.

A third limitation is external comparison. Our controlled BM25 and iterative lexical rows test retrieval behavior inside our candidate-pool protocol, but they are not official reproductions of LeanSearch v2 or other full premise-selection systems. The present verified result is therefore best read as orthogonal to global retrieval: once a retriever proposes names, Lean still needs a typed interface decision for how those names enter proof search.

8 CONCLUSION

This paper argues that Lean premise selection is incomplete without action-conditional evidence allocation. ACE implements this view as a typed evidence compiler over Hammer, HammerCore, simplification, `solve_by_elim`, and Aesop interfaces. On replayable Mathlib theorem contexts, the typed action grid reaches 58/230 verified proofs versus 38/230 for the best single action, and a four-retry fixed typed control recovers 57/230 out of fold. The exact Aesop controls make the final claim more precise: joint oracle-core plus retrieved exposure is the best Aesop row, but the margin over identity, simpsons-only, and single-source controls is modest. This positions typed evidence allocation as a concrete mechanism for building more reliable Lean hammers, with learned adaptive allocation and stronger external retriever-only validation remaining the next method targets rather than claims already supported by the current verified results.

REFERENCES

Guoxiong Gao, Zeming Sun, Jiedong Jiang, Yutong Wang, Jingda Xu, Peihao Wu, Bryan Dai, and Bin Dong. Leansearch v2: Global premise retrieval for lean 4 theorem proving. *arXiv preprint arXiv:2605.13137*, 2026. URL <https://arxiv.org/abs/2605.13137>.

- Thibault Gauthier, Cezary Kaliszyk, Josef Urban, Ramana Kumar, and Michael Norrish. Tactictoe: Learning to prove with tactics. *Journal of Automated Reasoning*, 65:257–286, 2021. URL <https://arxiv.org/abs/1804.00596>.
- Edward J. Hu et al. minictx: Neural theorem proving with long contexts. *arXiv preprint arXiv:2408.03350*, 2024. URL <https://arxiv.org/abs/2408.03350>.
- Suozhi Huang, Peiyang Song, Robert Joseph George, and Anima Anandkumar. Leanprogress: Guiding search for neural theorem proving via proof progress prediction. *arXiv preprint arXiv:2502.17925*, 2025. URL <https://arxiv.org/abs/2502.17925>.
- Minsu Kim and Se-Young Yun. Process-verified reinforcement learning for theorem proving via lean. In *International Conference on Learning Representations*, 2026. URL <https://openreview.net/forum?id=P00k4DFaXF>.
- David Ma, Kaijing Ma, Shawn Guo, Yunfeng Shi, Enduo Zhao, Jiajun Shi, Zhaoxiang Zhang, Gavin Cheung, Jiaheng Liu, and Zili Wang. Oprover: A unified framework for agentic formal theorem proving. *arXiv preprint arXiv:2605.17283*, 2026. URL <https://arxiv.org/abs/2605.17283>.
- Azim Ospanov, Farzan Farnia, and Roozbeh Yousefzadeh. Apollo: Automated llm and lean collaboration for advanced formal reasoning. In *Advances in Neural Information Processing Systems*, 2025. URL <https://openreview.net/forum?id=fxDCgOruk0>.
- Evan Wang, Simon Chess, Daniel Lee, Siyuan Ge, Ajit Mallavarapu, and Vasily Il'in. Learning to repair lean proofs from compiler feedback. *arXiv preprint arXiv:2602.02990*, 2026. URL <https://arxiv.org/abs/2602.02990>.
- Kaiyu Yang, Aidan M. Swope, Alex Gu, Rahul Chalamala, Peiyang Song, Shixing Yu, Saad Godil, Ryan Prenger, and Anima Anandkumar. Leandojo: Theorem proving with retrieval-augmented language models. In *Advances in Neural Information Processing Systems*, 2023. URL <https://arxiv.org/abs/2306.15626>.
- Thomas Zhu, Joshua Clune, Jeremy Avigad, Albert Qiaochu Jiang, and Sean Welleck. Premise selection for a lean hammer. *arXiv preprint arXiv:2506.07477*, 2025. URL <https://arxiv.org/abs/2506.07477>.

A ADDITIONAL EXPERIMENTAL DETAILS

A.1 CANONICAL OUTPUT SOURCES

All numbers in the main tables are copied from canonical JSON or Markdown outputs in the project directory:

- Verified typed action grid: `outputs/mathlib430_pretheorem_action_matrix_scaled230_aesop_ablation_merged.md`.
- Budgeted typed portfolio: `outputs/mathlib430_budgeted_action_policy_scaled230_aesop_ablation.md`.
- Typed allocator gate: `outputs/mathlib430_typed_allocator_gate_scaled230_aesop_ablation.md`.
- Exact Aesop source/channel controls: `analysis/mathlib430_aesop_exact_controls_summary.md`.
- Initial broad Aesop fact/simp ablation: `analysis/mathlib430_aesop_ablation_scaled230.md`.
- Initial Aesop counterfactual channel control: `analysis/mathlib430_aesop_counterfactual_controls.md`.
- Evidence-contract audit and homogeneous controls: `analysis/mathlib430_evidence_contract_audit.md`.
- Strict interface filtering boundary: `analysis/mathlib430_e1_strict_interface_filtering.md`.

- Strict interface taxonomy: `analysis/mathlib430_strict_action_taxonomy_scaled230_aesop_ablation.md`.
- Local trace-core: `outputs/phase1_mathlib_trace_core_2000_feature_ablation_a40.json`.
- Phase 1 bridge: `outputs/phase1_reconstruction_bridge_report_100.md`.
- Timeout stress: `outputs/phase2_timeout_stress_2000_a40.json`.
- Imported-core heldout: `outputs/phase3_second_stage_controller_eval_500_a40.json` plus final guardrail deltas.
- Split stability: `outputs/phase3_split_stability_summary_4x500.md`.
- Imported bridge report: `outputs/phase3_bridge_report_200.md`.
- Imported bridge taxonomy: `outputs/phase3_bridge_failure_taxonomy_200.md`.
- Imported bridge inspection: `outputs/phase3_bridge_negative_case_inspection_200.md`.

A.2 TRACE-CORE DISCOVERY SUMMARY

Table 8: Representative trace-core discovery results that motivated the verified typed-action protocol. These rows are not the main proving metric; they explain the design pressure for bounded retries, non-monotonic premise budgets, and bridge checks.

Discovery question	Strongest useful result	Readout for current paper
Can failure feedback improve local premise recovery?	Local trace-core FAR-HAMMER reaches 95.7% trace success and 72.0% bridge verified success; blind top- k expansion reaches 91.5% trace and 43.0% bridge.	Failure events can change which evidence is useful, but bridge verification is stricter than trace-core recovery.
Can adding more premises hurt?	Timeout-conditioned shrink reaches 69.3% success with 23.5% timeout rate, while equal-budget top- k falls to 43.8% with 56.2% timeouts.	Premise coverage is not monotonic under bounded search, motivating typed budgets instead of broad insertion.
Do local heuristics solve global imported premises?	Imported-core lexical/BM25 stress stays near 12% success before learned retrieval.	Local feature control is insufficient for global retrieval; the verified paper should not claim full retriever replacement.
Does learned imported retrieval survive bridge checks?	Final-base8 averages 92.9% across four splits and reaches 53.5% imported bridge verified success versus 44.5% for learned+base fallback.	Bridge evidence supports the premise-recovery story but remains separate from open-ended proof generation.

A.3 FEATURE-GROUP ABLATION

Table 9: Feature-group ablation on 2k local trace-core goals. Static visible-feature reranking is weaker than blind expansion; failure-conditioned feature use is the important step.

Method	Success	FFR	Avg prem.
<code>topk_expansion</code>	91.5%	79.9%	46.9
<code>visible_feature_name_rerank</code>	85.3%	0.0%	42.9
<code>visible_feature_statement_rerank</code>	77.6%	0.0%	42.9
<code>visible_feature_decl_rerank</code>	78.3%	0.0%	42.9
<code>visible_feature_rerank</code>	86.5%	0.0%	42.9
<code>rule_far_no_core_decl_features</code>	92.6%	82.4%	46.2
<code>rule_far_no_core_statement_features</code>	92.8%	82.9%	46.4
<code>rule_far_no_core_name_features</code>	95.0%	88.1%	45.4
<code>rule_far_no_core_tags</code>	95.7%	89.7%	45.1

A.4 IMPORTED-CORE STRESS BEFORE LEARNED RETRIEVAL

Table 10: Initial imported-core stress test on 2k goals. This table is a diagnostic negative result: local and lexical signals are insufficient for global imported-premise recovery.

Method	Success	FFR	Avg prem.
one_shot	1.1%	0.0%	32.0
same_file_prior_rerank	2.1%	0.0%	56.0
visible_feature_rerank	2.4%	0.0%	56.0
bm25_rerank	7.0%	0.0%	56.0
leansearch_iterative	11.9%	8.3%	91.9
bm25_expansion	11.9%	8.3%	92.2
rule_far_bm25	12.2%	8.7%	92.2
rule_far_full (oracle ref.)	55.7%	55.2%	86.6

A.5 BRIDGE FAILURE TAXONOMY

The 200-goal imported bridge verifies 124/200 original traced tactic scripts. The primary replay failures are 30 unknown-identifier errors, 13 other Lean errors, 8 typeclass failures, 7 unsolved-goal failures, 6 simp no-progress failures, 5 tactic failures, 5 type mismatches, and 2 sorry-context failures. At the policy level, final-base8 has 17 replay-verified trace misses, compared with 19 for the original second-stage controller, 35 for learned+base fallback, and 47 for learned expansion. Negative-case inspection finds 0/17 candidate-pool misses, 2 fallback-selected core drops, 3 expansion-selected core drops, and 15 failure-conditioned rank misses. This supports the limitation that remaining headroom is ranking and calibration, not candidate generation.

A.6 GUARDRAIL AND EXPERT-MIXING ABLATIONS

The final base8 guardrail was selected because broader heuristics were negative. Hard base guardrails reduce bridge trace-core success from 72.0% to 63.0% or 60.0%. Soft base-score mixing is inert at low mixing weights and harmful at high weights. Imported-failure multi-expert routing drops bridge trace-core success to 62.0%. A larger final expert guardrail also underperforms final-base8 on the original split, fold0, and bridge subset. These results justify stopping hand-written expert-mix tuning and treating calibrated expert gating as future work.

B CLAIM-EVIDENCE SELF-REVIEW

Table 11: End-of-draft self-review. Claims in the abstract and introduction are aligned to reported evidence; unsupported stronger claims are explicitly avoided.

Dimension	Current answer	Residual risk
Contribution	Action-conditional evidence allocation is formulated as distinct from premise ranking, with Aesop channel counterfactuals as mechanism evidence.	A fixed portfolio can look simple unless the interface mechanism remains central.
Writing clarity	Main text now leads with verified typed actions, then uses trace-core as discovery evidence.	Method names are long; final draft may need shorter aliases.
Experiments	Results cover verified replayable Mathlib actions, budgeted allocators, Aesop counterfactuals, trace-core discovery, and replay bridges.	Replayable subset is 230/490, not all traces.
Evaluation	Kernel-verified action success, oracle headroom, fixed-budget policy, and interface ablations are reported.	Full open-ended theorem proving is out of scope.
Soundness	Negative gates show broad insertion, strict syntactic filtering, and learned adaptive allocation are not enough.	Learned allocation does not yet beat the fixed typed control.